

Parallel Computing Systems (CS4171-KP12)

Exercise Sheet 1 Submission

Md Norul Gofran

`md.gofran@student.uni-luebeck.de`
Matriculation Number: 814021

November 3, 2025

1 Task 1: Think Parallel (Prefix Sum)

Given Array (N=8): $x = [4, 9, 2, 7, 1, 6, 3, 8]$

1.1 Task 1.1: Sequential and Parallel Algorithm Design

1.1.1 Sequential Algorithm

The sequential algorithm for an inclusive prefix sum:

```
y[0] = x[0]
FOR i = 1 to N-1
    y[i] = y[i-1] + x[i]
END FOR
```

Execution Trace:

- $y[0] = 4$
- $y[1] = 4 + 9 = 13$
- $y[2] = 13 + 2 = 15$
- $y[3] = 15 + 7 = 22$
- $y[4] = 22 + 1 = 23$
- $y[5] = 23 + 6 = 29$
- $y[6] = 29 + 3 = 32$
- $y[7] = 32 + 8 = 40$

Final Prefix Sum: $[4, 13, 15, 22, 23, 29, 32, 40]$

1.1.2 Parallel Algorithm (Blelloch Scan)

The parallel algorithm consists of two passes (up-sweep and down-sweep).

1. **Reduce (Up-Sweep) Phase:** Builds a binary tree of partial sums.
2. **Down-Sweep Phase:** Traverses the tree from the root down to build the exclusive prefix sum.
3. **Final Step (Exclusive to Inclusive):** A final parallel operation where $y[i] = y_{\text{exclusive}}[i] + x[i]$.

1.2 Task 1.2: Reduce Operation (Up-Sweep Phase)

Input: [4, 9, 2, 7, 1, 6, 3, 8]

| Stride | Active Indices (i) | Array State x |
|---------|--------------------|-----------------------------|
| Initial | - | [4, 9, 2, 7, 1, 6, 3, 8] |
| 1 | 1, 3, 5, 7 | [4, 13, 2, 9, 1, 7, 3, 11] |
| 2 | 3, 7 | [4, 13, 2, 22, 1, 7, 3, 18] |
| 4 | 7 | [4, 13, 2, 22, 1, 7, 3, 40] |

Outcome of Reduce Operation: The total sum, **40**, is in $x[7]$.

1.3 Task 1.3: Down-Sweep Phase Output

This phase constructs the *exclusive* prefix sum. We set the last element to 0.

Start: [4, 13, 2, 22, 1, 7, 3, 40] \rightarrow [4, 13, 2, 22, 1, 7, 3, 0]

| Stride | Active Indices (i) | Array State x |
|--------|--------------------|--------------------------------|
| 4 | 3, 7 | [4, 13, 2, 0, 1, 7, 3, 22] |
| 2 | 1, 3, 5, 7 | [4, 0, 2, 13, 1, 22, 3, 29] |
| 1 | 1, 3, 5, 7 | [0, 4, 13, 15, 22, 23, 29, 32] |

Final Output of Down-Sweep (Exclusive Sum): [0, 4, 13, 15, 22, 23, 29, 32]

1.4 Task 1.4: Evaluation and Complexity

| Metric | Sequential Strategy | Parallel Strategy (Blelloch) |
|------------------------|-------------------------------------|---|
| Time | $O(N)$ (One loop of $N-1$ steps) | $O(\log N)$ (Two tree traversals of $\log N$ steps) |
| Operation Count | $O(N)$ ($N - 1$ additions) | $O(N)$ ($2(N - 1)$ additions for up/down) |
| Required CPUs | 1 (Inherently serial) | $O(N)$ (specifically $N/2$) (To achieve $O(\log N)$ time) |

2 Task 2: Parallel Algorithms (Scalar Product)

Given: Two vectors, A and B, with $N = 200$ elements. $p = 8$ processors.

2.1 Task 2.1: Sequential Algorithm

```
sum = 0
FOR i = 0 to N-1
    product = A[i] * B[i]
    sum = sum + product
END FOR
RETURN sum
```

2.2 Task 2.2: Parallel Algorithm (8 Processors)

1. **Distribute Work:** Each processor is responsible for $N/p = 200/8 = 25$ elements.
2. **Local Computation (Parallel):** Each processor k calculates its own local scalar product (a partial sum) for its 25 elements.
3. **Global Reduction (Parallel):** The 8 partial sums are combined into a final sum using a parallel reduction.

Scheme:

```
// --- Phase 1 & 2: Local Computation (in parallel by all 8) ---
// Processor 'k' (where k is 0 to 7)
local_sum_k = 0
my_start = k * 25
my_end = my_start + 25
FOR i = my_start to my_end-1
    local_sum_k = local_sum_k + (A[i] * B[i])
END FOR

// --- Phase 3: Global Reduction (in log_2(p) steps) ---
global_sum = REDUCE(local_sum_k, operation=SUM)
RETURN global_sum
```

2.3 Task 2.3: Time Steps

(Assuming 1 time step per operation)

- **Sequential Time (T_1):**
 - N multiplications + $N - 1$ additions
 - $T_1 = 200 + 199 = \mathbf{399}$ time steps.
- **Parallel Time (T_8):**
 - **Local Compute:** (Done by all processors in parallel)
 - $N/p = 25$ multiplications
 - $(N/p - 1) = 24$ additions
 - Time = $25 + 24 = 49$ steps.
 - **Global Reduction:** Add 8 partial sums.
 - Time = $\log_2(p) = \log_2(8) = 3$ addition steps.
 - $T_8 = 49 + 3 = \mathbf{52}$ time steps.

2.4 Task 2.4: Speed-up and Efficiency ($p = 8$)

- **Speed-up (S_8):** $S_p = T_1/T_p$

$$S_8 = \frac{399}{52} \approx \mathbf{7.67}$$

- **Efficiency (E_8):** $E_p = S_p/p$

$$E_8 = \frac{7.67}{8} \approx 0.959 \text{ (or } \mathbf{95.9\%})$$

2.5 Task 2.5: Generalized Functions (N and p)

- **Time (Sequential):**

$$T_1(n) = n + (n - 1) = 2n - 1$$

- **Time (Parallel):**

$$T_p(n, p) = \underbrace{\left(\frac{n}{p} + \left(\frac{n}{p} - 1 \right) \right)}_{\text{Local Compute}} + \underbrace{\log_2 p}_{\text{Global Reduce}}$$

$$T_p(n, p) = \frac{2n}{p} - 1 + \log_2 p$$

- **Speed-up (S_p):**

$$S_p(n, p) = \frac{T_1}{T_p} = \frac{2n - 1}{\frac{2n}{p} - 1 + \log_2 p}$$

- **Efficiency (E_p):**

$$E_p(n, p) = \frac{S_p}{p} = \frac{2n - 1}{p \left(\frac{2n}{p} - 1 + \log_2 p \right)} = \frac{2n - 1}{2n - p + p \log_2 p}$$

3 Task 3: PRAM - Matrix Multiplication

Given $N \times N$ matrices A, B, and C. Sequential time $T_1 = O(n^3)$.

3.1 Task 3.1: EREW PRAM Algorithm

An EREW algorithm requires explicit broadcast phases.

Algorithm (using $p = n^3$ processors $P_{i,j,k}$):

1. Phase 1: Broadcast Data (EREW)

- In parallel: Broadcast $A[i, k]$ to all n processors $P_{i,j,k}$ (for $j = 1..n$).
- In parallel: Broadcast $B[k, j]$ to all n processors $P_{i,j,k}$ (for $i = 1..n$).
- EREW broadcast takes $O(\log n)$ time.

2. Phase 2: Compute Products (Parallel)

- In parallel, all n^3 processors $P_{i,j,k}$ compute:
- $\text{Temp}[i, j, k] = A[i, k] * B[k, j]$

3. Phase 3: Reduce Sums (EREW)

- In parallel for all i, j : Perform a parallel sum (reduction) of the n values $\text{Temp}[i, j, k]$ (for $k = 1..n$).
- EREW summation takes $O(\log n)$ time.
- $C[i, j] = \sum_{k=1}^n \text{Temp}[i, j, k]$

- **Processors Used:** $O(n^3)$

- **Time Complexity:** $O(\log n) + O(1) + O(\log n) = \Theta(\log n)$.

3.2 Task 3.2: $O(n^2)$ Processors Algorithm (CREW)

We assign one processor $P_{i,j}$ to compute one element $C[i,j]$. This requires a **CREW PRAM model** because multiple processors read the same rows of A and columns of B.

Algorithm (using $p = n^2$ processors):

1. Phase 1: Compute Products (Sequential per Processor)

- In parallel, each processor $P_{i,j}$ does:
- FOR $k = 1$ to n :
- $\text{Temp}[k] = A[i,k] * B[k,j] // Concurrent Read$

2. Phase 2: Compute Sum (Sequential per Processor)

- In parallel, each processor $P_{i,j}$ does:
- $C[i,j] = 0$
- FOR $k = 1$ to n :
- $C[i,j] = C[i,j] + \text{Temp}[k] // Exclusive Write$
- **Time Complexity:** $T_p = O(n) + O(n) = \mathbf{O}(n)$.

• Speed-up (S_p):

$$S_p = \frac{T_1}{T_p} = \frac{O(n^3)}{O(n)} = \mathbf{O}(n^2)$$

• Efficiency (E_p):

$$E_p = \frac{S_p}{p} = \frac{\mathbf{O}(n^2)}{O(n^2)} = \mathbf{O}(1)$$

3.2.1 Comparison

The $O(n^2)$ processor algorithm is slower ($O(n)$) than the $O(n^3)$ processor algorithm ($O(\log n)$), but it is **work-optimal** and has an efficiency of $O(1)$, making it a more practical use of resources.

4 Task 4: Distributed Search

Given: A string of length n with unique characters. p processors. Find index of 'c'.

4.1 Task 4.1: Parallel Search Algorithms

All models share the same first step:

1. Phase 1: Local Search (All Models)

- Divide the string into p chunks of size n/p .
- In parallel, each processor k sequentially searches its local chunk.
- This takes $O(n/p)$ time (worst case).
- Each processor k stores its result: `local_index = (found_index or -1)`.

2. Phase 2: Global Reduction (Model-Dependent)

This phase finds the single non-negative index from the p `local_index` variables. A `global_result` is initialized to -1.

4.1.1 EREW-PRAM Algorithm

- **Phase 2 (EREW Reduction):** A reduction tree is used to find the one valid index. In $\log_2 p$ steps, processors compare results in pairs and pass up the non-negative index.

4.1.2 CREW-PRAM Algorithm

- **Phase 2 (CREW Reduction):** The $O(\log p)$ EREW reduction algorithm is the most efficient method and works on a CREW machine.

4.1.3 CRCW-PRAM Algorithm

- **Phase 2 (CRCW Reduction):** We use the **Arbitrary Write** CRCW model.
 - If processor k finds the character (`local_index` $\neq -1$), it writes its result to the shared `global_result`.
 - IF `local_index` $\neq -1$ THEN `global_result = i`
 - Since the character is unique, only one processor writes, completing in $O(1)$ time.

4.2 Task 4.2: Time Complexity Analysis

- **EREW-PRAM:**

$$T_p = \underbrace{O(n/p)}_{\text{Local Search}} + \underbrace{O(\log p)}_{\text{EREW Reduce}} = \mathbf{O}(n/p + \log p)$$

- **CREW-PRAM:**

$$T_p = \underbrace{O(n/p)}_{\text{Local Search}} + \underbrace{O(\log p)}_{\text{Best Reduce}} = \mathbf{O}(n/p + \log p)$$

- **CRCW-PRAM (Arbitrary):**

$$T_p = \underbrace{O(n/p)}_{\text{Local Search}} + \underbrace{O(1)}_{\text{CRCW Write}} = \mathbf{O}(n/p)$$